

Parallisation concept

Fatema Alkhanaizi

January 17, 2019

The first loop in `updateBody()` is for computing the force. The force for each particle is the sum of all forces acting on it from the rest of the particles in the system; a nested loop is used to implement this:

```
1  for (i = 0; i < NumberOfBodies; ++i) {
2      xi = x[i][0];
3      yi = x[i][1];
4      zi = x[i][2];
5      fx = 0.0;
6      fy = 0.0;
7      fz = 0.0;
8
9      // compute force on particle i
10     for (j = 0; j < NumberOfBodies; j++) {
11         ...
12         minDx = std::min(minDx, r2);
13     }
14
15     force[i][0] = fx;
16     force[i][1] = fy;
17     force[i][2] = fz;
18 }
```

Each iteration of the outer loop `i` is independent and the inner loops deals with computing the value for the force vector for the `i`th particle. There are a lot of duplicated computations, e.g. the force when particle `i=0` and particle `j=1` is the same as the force when particle `i=1` and particle `j=0` in terms of magnitude, so it is possible to evaluate them in one step. However, this could cause a race condition in which if running the nested loop in parallel two or more threads will try to modify the same particle at the same time. If a critical section is introduced this will slow the computation for the threads which could give worst results than if ran in serial due to the overhead introduced by parallization in OpenMP. To reduce the access to the force array, variables were setup to compute the force vector. The minimum distance is computed accross all iterations in the both loops. It could be determined in a reduction steps following Bulk Synchronous programming (BSP) method which OpenMP allows; it will be reduction for the minimum value. As OpenMP performs on a shared memory, any variable initialized outside the parallel region will be shared so the scalar variables inside the loop that will be constantly modified will need to be set as private.

The second loop in `updateBody()` deals with updating the movement of the particles:

```

1  for (i = 0; i < NumberOfBodies; i++) {
2
3      x[i][0] = x[i][0] + timeStepSize * v[i][0];
4      x[i][1] = x[i][1] + timeStepSize * v[i][1];
5      x[i][2] = x[i][2] + timeStepSize * v[i][2];
6
7      mt = timeStepSize / mass[i];
8
9      v[i][0] = v[i][0] + mt * force[i][0];
10     v[i][1] = v[i][1] + mt * force[i][1];
11     v[i][2] = v[i][2] + mt * force[i][2];
12
13     ...
14
15     maxV = std::max(maxV, V);
16 }

```

Each iteration `i` is independent in terms of access to `x` and `v` arrays except for computing the `maxV`. For `maxV`, a reduction step with OpenMP is possible as `maxV` is a scalar (reduction for `max`), thus it is possible to run this loop concurrently. Similarly the scalar variables inside the loop will need to be set to private; similarly to the previous loop.

For measuring performance, relative speed and efficiency measurement will be used:

$$S(p) = \frac{t(1)}{t(p)}$$

$$E(p) = \frac{S(p)}{p}$$

By following Amdahl's law, an increase in speed and efficiency is expected when the number of cores is increased while the number of bodies remains constant.